

# Solving a McCall Search Model via Reinforcement Learning

Trevor Gallen

Econ 64200

Fall 2022

## Deliverables

- You should have a word/L<sup>A</sup>T<sub>E</sub>X document that has three sections:
  1. Discusses the model and answers the questions I pose throughout.
  2. Contains the tables and figures you will produce.
  3. Contains a discussion of your programming choices if you had to make any.
- You should have a Matlab file or set of files (zipped) that contain **all** your programs and raw data. There should be a file called “Main.M” that produces everything I need in one click.

## 1 Model

In this homework, we’re going to try to solve a simple McCall search model with Reinforcement Learning.

Each period, agents wage up with a wage draw  $w \sim \log \mathcal{N}(1, 3)$ , they have a choice whether or not to accept the offer or not. If they accept, they receive that  $w^{accepted}$  draw every period for the rest of their life. If they reject, then (if they don’t die) receive zero and get a new wage draw in the next period, discounted by a factor of  $\beta$ . The probability of death is  $\alpha = 0.067$ . Their problem is can therefore be written:

$$V(w) = \max_{accept, reject} \left\{ \frac{w}{1 - (1 - \alpha)\beta}, (1 - \alpha)\beta V(w') \right\}$$

## 2 Problem

Set this problem up and solve it with a reinforcement learning agent in Matlab. Note that while the state is continuous, the decision is discrete. However, you may make a discrete decision continuous (round a sigmoid function, for instance), and approximate a continuous state with a discrete set of  $N$  levels. To solve this, you must therefore:

1. Set up an action space (accept/reject) and observation space (wage draw)
2. Set up neural networks (if actor critic, then a critic network that takes in one action and one state and spits out a scalar, and an action network that takes in a state and spits out an action)
3. Set up a reset function that starts your actor
4. Set up a step function that advances them one period, or terminates the problem if they accept (or die)

This can be as simple as creating the observation info (`rlNumericSpec(dimensions)`), the act info (`rlFiniteSetSpec(discrete actions)`), set up the environment (`rlFunctionEnv(obsInfo,actInfo,stepfcn,resetfcn)`) and then setting up a default DQN agent (`rlDQNAgent(obsInfo,actInfo)`), telling it the discount factor (`agent.AgentOptions.DiscountFactor = xxx`), and training it `train(agent,env,opt)`. **I leave it to you** how you solve it, this is merely a suggestion.

Once you've solved the problem, you should graph out:

- The value function of each action, contingent on draw (`getValue(getCritic(agent,wage draw))`)
- The action at each wage draw `getAction(agent,wage draw)`

## 3 Research

Change the model in some interesting way and report your results.