

# NUMERICAL METHODS-LECTURE IX: REINFORCEMENT LEARNING

(See Powell Chapter 10)

Trevor S. Gallen

# MOTIVATION

- ▶ Previously, this lecture was devoted to quadrature, and specifically how to integrate sampling from Chebyshev nodes
- ▶ While valuable, it may be going out of style relative to modern methods
- ▶ Instead, we'll talk about Reinforcement Learning
- ▶ Could have had a whole course, but I want to give you the motivation and a crash course

# THE THREE CURSES OF DIMENSIONALITY

► Recall the three curses of dimensionality

1. As  $n_{\text{states}}$  proliferates, # problems to solve explodes for VFI
2. As  $n_{\text{actions}}$  proliferates, # choices to compare explodes
3. As  $n_{\text{outcomes}}$  proliferates (particularly stochastic), space to integrate over explodes

$$V(x) = \max_{y \in \Gamma(X)} \{F(x, y) + \beta E(V(x'(y)))\}$$

- This is a problem if you're trying to model lifecycle behavior...age, permanent wage, transitory wage, marital status, age of kids, occupation, health, etc.
- How can we solve?

# ONE WAY TO APPROXIMATELY SOLVE THE PROBLEM

- ▶ There are many flavors of solution, but we'll focus on the "Actor-Critic" method
- ▶ Have two functions, parameterized by  $\theta$  and  $\phi$ :
  - ▶ Actor function  $\bar{\pi}(y|x;\theta)$  takes in state and spits out an action (possibly probabilistically)
  - ▶ Critic function  $\bar{V}(x|\phi)$  takes in state and spits out value (traditional value function)
- ▶ Good actor function embeds both reward and stochastic future (actions and integral!)
- ▶ *Given*  $\theta$  and  $\phi$ , can simulate an agent
- ▶ Need to find  $\theta$  and  $\phi$

# ACTOR-CRITIC ALGORITHM (DISCUSSION)

- ▶ Start with a guess for  $\theta$  and  $\phi$ , and an initial state value
- ▶ Simulate the system many times (random draws & laws of motion for stochastic problems)
- ▶ Now we have a bunch of paths for a given  $\theta$  and  $\phi$
- ▶ For every step  $t$ , compute the return  $G_t$ , sum of reward and discounted future reward, calculated with  $\overline{V(x_{t+1})}$
- ▶ Calculate the “advantage function”  $D_t = G_t - V(S_t|\phi)$ , value of action vs value of what we think is best action embedded in  $V$
- ▶ Calculate gradients for actor and critic networks:

$$d\theta = \sum_{i=1}^N \nabla_{\theta} \log(\pi(A|x_t; \theta)) D_t$$

$$d\phi = \sum_{t=1}^N \nabla_{\phi} (G_t - V(x_t; \phi))^2$$

- ▶ Idea:  $\nabla_{\theta}$  pushes us in direction of better choices/happiness,  $\nabla_{\phi}$  pushes us in direction of lower error in value function
- ▶ Update the parameters of the functions:

$$\theta = \theta + \alpha d\theta$$

$$\phi = \phi + \beta d\phi$$

- ▶ So we update the actor  $\pi$  using critic  $V$ , and update  $V$  using simulation

## EXPLAINING TO MUM-I

- ▶ Start out with some surface that represents best action given state, and some surface that represents the value of being in that state
- ▶ Simulate
- ▶ Change value surface by comparing data with what you actually got (using initial surface for future, so change is slow), trying to match surface
- ▶ Change action surface by trying to increase received value vs value at best guess (small perturbations toward better actions)
- ▶ Repeat in tiny steps

# IDEA

- ▶ We try to toddle slowly to both how to evaluate our situation ( $V$ ) and what to do ( $\pi$ )
- ▶ We learn about value function by exploring the space
- ▶ We learn about maximization (actor) by exploiting  $V$  and our actual actions
- ▶ We learn about expectation by simulation—enough simulations and we will explore the relevant space
- ▶ Additional cool (but dangerous) aspect: we only explore relevant functions of state space
- ▶ But how choose  $V$  and  $\pi$ ?

## CHOOSING $V$ AND $\pi$

- ▶  $V$  and  $\pi$  could be any functional form (e.g. linear  $V = \alpha + \beta \sum_{j=1}^N \phi_j x_j$ )
- ▶ But we want an extremely functional flexible form
- ▶ Neural networks are (typically) just simple stacked functions interacted with one another—think very flexible functions, with many parameters
- ▶ Advantage of NN-style flexibility is we can spend degrees of freedom in complicated areas and not in simple areas (as in sparse interpolation)
- ▶ I'm giving this short shrift (sorry), but we'll put together a neural network for  $V$  and  $\pi$  in Matlab



# PROBLEM TO SOLVE

- ▶ We'll solve a simple finite-horizon NCG-style problem:

$$V(A, K, t) = \max_{K'} \{ \log(AK^\alpha - K') + \beta E(V(A', K', t+1)) \}$$

$$A' = (1 - \rho) + \rho A + \epsilon, \quad \epsilon \sim \mathcal{N}(0, 0.05)$$

- ▶ How will we set this up?
  - ▶ Define a set of functions that:
    - ▶ A function that sets up actor and critic neural networks, defines observations, and trains (Main.m)
    - ▶ Initialize an agent (including random draws) (myResetFunction.m)
    - ▶ Step an agent forward in time, simulate draws (myStepFunction.m)

## RESET FUNCTION

```
function [InitialObservation, LoggedSignal] =  
myResetFunction()  
    % Initial values  
    S.K = 250+(rand(1,1)-0.5)*200;  
    S.A = 1+rand(1,1)*0.05;  
    S.step = 1  
    % Return initial environment state variables as  
    logged signals.  
    LoggedSignal.State = [S.K;S.A;S.step];  
    LoggedSignal.C=NaN;  
    InitialObservation = LoggedSignal.State;
```

## STEP FUNCTION

- ▶ See myStepFunction.m
- ▶ This is the simulator, it takes in signals and an action and simulates the environment, returns the observations, rewards, and whether or not it is done
- ▶ Note: logged signals are known to Matlab, but not agent. Observations are known to agent. In this simulation, they are the same thing. (But could have had hidden state)

# MAIN.M

- ▶ See main.M
- ▶ Idea: set up a flexible function for the actor and critic, and then send to trainer


## COMPARING TO TRADITIONAL VFI

- ▶ There are many options to maximize, can take more time, etc. but let's get a ballpark idea of how well this can do
- ▶ Hard to see, so let's look at GraphDiff.m

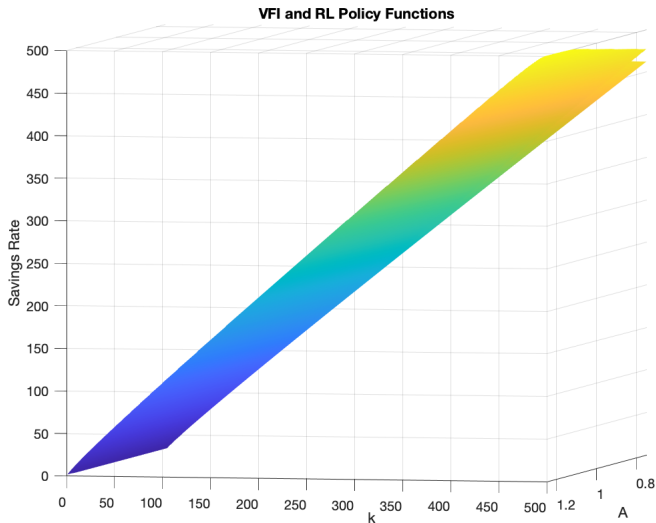
## VALUE FUNCTIONS ARE SIMILAR

- ▶ There are many options to maximize, can take more time, etc. but let's get a ballpark idea of how well this can do after five hours
- ▶ Note: won't be perfect! Feel free to run longer and/or with more parameters
- ▶ Hard to see, so let's look at GraphDiff.m

# VALUE FUNCTIONS



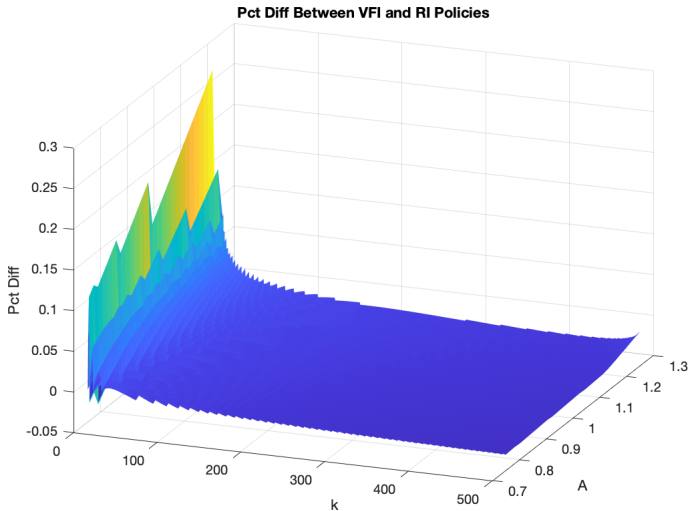
# POLICY FUNCTIONS



Hard to see, but right on top of one another

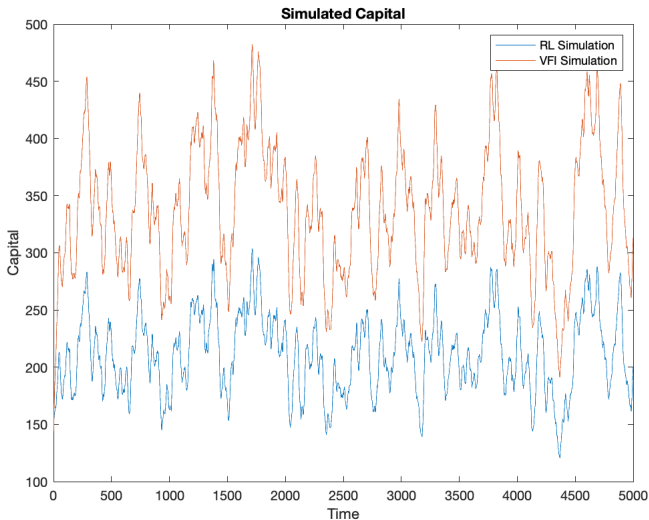


# POLICY FUNCTIONS DIFFERENCE




Mostly very small differences...except at  $k$  near zero (why?)

# SIMULATION



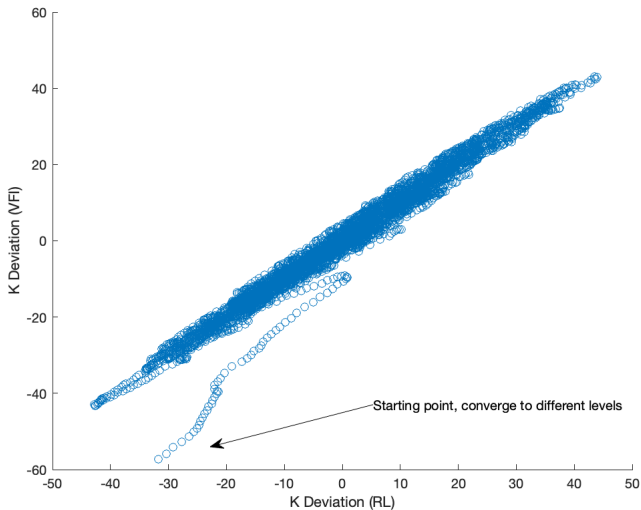
We get towards the right answer in percent deviations (let's check!)

# SIMULATION



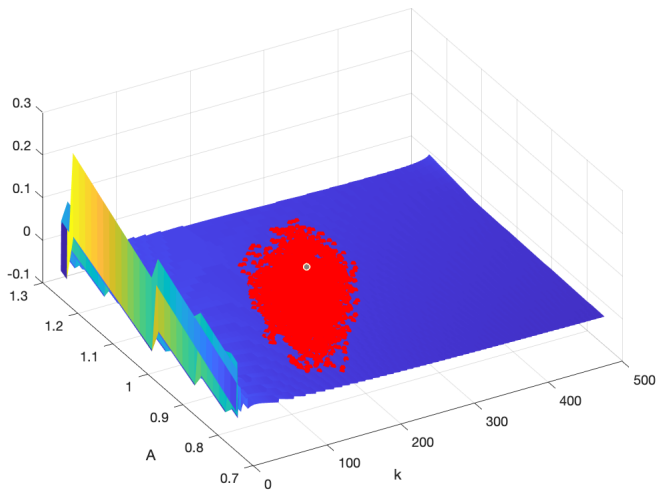
Mostly spot-on in differences

# SIMULATION




Not terrible, but could use more running & debugging to get level

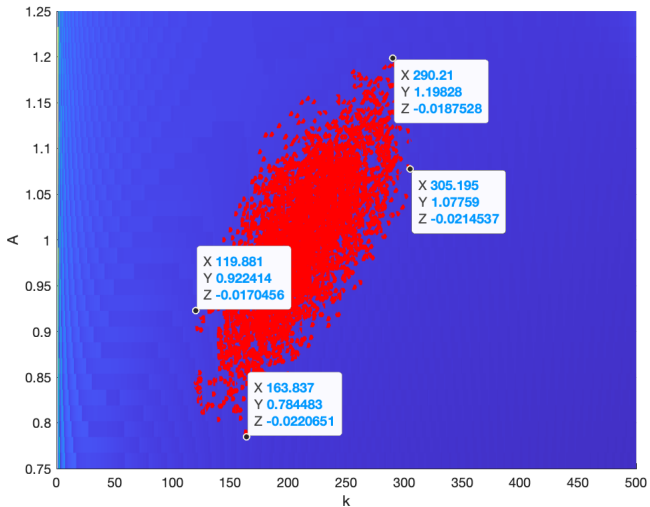
WE MOSTLY SAMPLE POINTS IN THE RED AREA, SO  
MOST ACCURATE THERE



WE MOSTLY SAMPLE POINTS IN THE RED AREA, SO  
MOST ACCURATE THERE



# SOME USEFUL BOUNDS



## SUMMING UP

- ▶ Reinforcement learning incredibly useful
- ▶ Nothing spectacularly clever, just a combination of:
  - ▶ Flexible functional forms (so many state variables can be accommodated efficiently)
  - ▶ *Forward-looking* (average over many simulations is estimation)
  - ▶ Simple maximization (use gradient and flexible function, so don't have to solve perfectly and solutions potentially informative across state space)
- ▶ Could have gone much farther. Could have had parameters  $\alpha$ ,  $\rho$ , etc. be draws too (solve not only over  $k$  and  $A$  but over parameterization, so can solve for heterogeneous agents or estimate parameters easily! (VFI would require solving for each parameter set).
- ▶ In practical terms, you just have a setup function, and then a function that steps through time (simulates) and pass it off to solver



## LAST WARNING

- ▶ There's a lifetime of details in terms of efficiency, problem setup, solvers, etc., and the devil is in the details
  - ▶ Discrete, continuous
  - ▶ Shallow RL, deep RL (how setup?)
  - ▶ On/off policy
  - ▶ Delayed learning
- ▶ We went through one algorithm (actor-critic) and one example (continuous state & action space).
- ▶ There are a cornucopia of flavors & algorithms—I haven't yet found one that isn't intuitive & obvious in retrospect (once you grok the core idea behind approximate dynamic programming idea)
- ▶ However, economist notation and ML notation diverge a bit, so there's an investment in learning